

Model Drift — Advanced

1 Label Drift (Target Drift)

What is Label Drift?

Label Drift occurs when the **distribution of target labels changes over time**, even if input features remain stable.

$$P(Y_{train}) \neq P(Y_{prod})$$

This is **the most dangerous drift** because:

- Feature drift tests won't detect it
- Model confidence can stay high
- Business metrics silently degrade

Why Label Drift Happens

- Economic changes (recession → more loan defaults)
- User behavior changes
- Policy or regulation changes
- Feedback loops caused by the model itself

Key Property (IMPORTANT)

⚠ Labels are delayed

You **cannot detect label drift instantly**.

You must wait for ground truth.

This is why:

- Feature drift = real-time
 - Label drift = lagged monitoring
-

How Label Drift Is Detected (Practically)

You do **NOT** compare distributions directly.

Instead, you monitor:

- Accuracy / F1 / AUC over time
- Error rate trend
- Calibration decay

Using **rolling windows**.

Python: Label Drift via Performance Decay

Example: Binary Classification

```
import numpy as np
from sklearn.metrics import accuracy_score, f1_score

# Simulated predictions and true labels
np.random.seed(42)

y_true_past = np.random.binomial(1, 0.3, 1000)
y_pred_past = y_true_past.copy() # good model

y_true_recent = np.random.binomial(1, 0.5, 1000) # label shift
y_pred_recent = np.random.binomial(1, 0.3, 1000) # model unchanged

# Metrics
past_f1 = f1_score(y_true_past, y_pred_past)
```

```

recent_f1 = f1_score(y_true_recent, y_pred_recent)

print(f"Past F1: {past_f1:.3f}")
print(f"Recent F1: {recent_f1:.3f}")

if recent_f1 < past_f1 * 0.9:
    print("⚠ LABEL DRIFT DETECTED (Performance Decay)")
else:
    print("✅ Performance Stable")

```

When to Act

Scenario	Action
Metric drop < 5%	Monitor
5–15%	Retrain
> 15%	Redesign model / features

2 Multivariate Data Drift

What is Multivariate Drift?

Multivariate drift occurs when **relationships between features change**, even if individual feature distributions look stable.

$$P(X_1, X_2, \dots, X_n) \neq P'(X_1, X_2, \dots, X_n)$$

Univariate tests (KS) **fail here**.

Real Example

- Age distribution: same

- Salary distribution: same
- **Age–Salary correlation changed**

Model breaks.

Technique: PCA Reconstruction Error

How PCA-Based Drift Detection Works

1. Fit PCA on training data
2. Reconstruct data using limited components
3. Measure reconstruction error
4. Compare baseline vs production error

If error increases → drift

Python: Multivariate Drift with PCA

```
import numpy as np
from sklearn.decomposition import PCA

np.random.seed(42)

# Training data (correlated features)
X_train = np.random.multivariate_normal(
    mean=[0,0],
    cov=[[1,0.9],
          [0.9,1]],
    size=1000
)

# Production data (correlation changed)
X_prod = np.random.multivariate_normal(
    mean=[0,0],
```

```

cov=[[1,0.2],
     [0.2,1]],
size=1000
)

# Fit PCA on training data
pca = PCA(n_components=1)
pca.fit(X_train)

# Reconstruction
X_train_recon = pca.inverse_transform(pca.transform(X_train))
X_prod_recon = pca.inverse_transform(pca.transform(X_prod))

# Reconstruction error
train_error = np.mean(np.linalg.norm(X_train - X_train_recon, axis=1))
prod_error = np.mean(np.linalg.norm(X_prod - X_prod_recon, axis=1))

print(f"Train Error: {train_error:.4f}")
print(f"Prod Error: {prod_error:.4f}")

if prod_error > train_error *1.2:
    print("⚠ MULTIVARIATE DRIFT DETECTED")
else:
    print("✅ No Multivariate Drift")

```

Why This Works

- PCA captures **feature interaction structure**
- Reconstruction error increases when relationships break
- Scales to many features